

Week 4 Exercises

Adeline Casali

July 28, 2023

Please complete all exercises below. You may use any library that we have covered in class. The data we will be using comes from the `tidyr` package, so you must use that.

- 1) Examine the `who` and `population` data sets that come with the `tidyr` library. the `who` data is not tidy, you will need to reshape the `new_sp_m014` to `newrel_f65` columns to long format retaining country, `iso2`, `iso3`, and year. The data in the columns you are reshaping contains patterns described in the details section below. You will need to assign three columns: diagnosis, gender, and age to the patterns described in the details.

Your tidy data should look like the following: country iso2 iso3 year diagnosis gender age count
1 Afghanistan AF AFG 1980 sp m 014 NA 2 Afghanistan AF AFG 1980 sp m 1524 NA 3 Afghanistan AF AFG 1980 sp m 2534 NA 4 Afghanistan AF AFG 1980 sp m 3544 NA 5 Afghanistan AF AFG 1980 sp m 4554 NA 6 Afghanistan AF AFG 1980 sp m 5564 NA

Details The data uses the original codes given by the World Health Organization. The column names for columns five through 60 are made by combining `new_` to a code for method of diagnosis (`rel` = relapse, `sn` = negative pulmonary smear, `sp` = positive pulmonary smear, `ep` = extrapulmonary) to a code for gender (`f` = female, `m` = male) to a code for age group (014 = 0-14 yrs of age, 1524 = 15-24 years of age, 2534 = 25 to 34 years of age, 3544 = 35 to 44 years of age, 4554 = 45 to 54 years of age, 5564 = 55 to 64 years of age, 65 = 65 years of age or older).

Note: use `data(who)` and `data(population)` to load the data into your environment. Use the arguments `cols`, `names_to`, `names_pattern`, and `values_to`. Your regex should be = (“new_?(.)_(.)(.)”)

<https://tidyr.tidyverse.org/reference/who.html>

```
# Load packages
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.2      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2    3.4.2      v tibble    3.2.1
## v lubridate  1.9.2      v tidyr     1.3.0
## v purrr      1.0.1
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
# Load data
data(who)
data(population)

# Convert wide to long format
tidy_who <- who %>%
  pivot_longer(cols = matches("new_?(.*)_(.)(.*)"),
               names_to = c("diagnosis", "gender", "age"),
               names_pattern = "new_?(.*)_(.)(.*)",
               values_to = "count") %>%
  mutate(age = ifelse(age == "65", "65", age))

head(tidy_who)
```

```
## # A tibble: 6 x 8
##   country    iso2 iso3   year diagnosis gender age   count
##   <chr>      <chr> <chr> <dbl> <chr>      <chr> <chr> <dbl>
## 1 Afghanistan AF    AFG   1980 sp        m    014    NA
## 2 Afghanistan AF    AFG   1980 sp        m   1524    NA
## 3 Afghanistan AF    AFG   1980 sp        m   2534    NA
## 4 Afghanistan AF    AFG   1980 sp        m   3544    NA
## 5 Afghanistan AF    AFG   1980 sp        m   4554    NA
## 6 Afghanistan AF    AFG   1980 sp        m   5564    NA
```

- 2) There are two common keys between the data sets, with who as the left table, join the population data by country and year so that the population is available within the who dataset.

```
tidy_who_pop <- left_join(tidy_who, population, by = c("country", "year"))

head(tidy_who_pop)
```

```
## # A tibble: 6 x 9
##   country    iso2 iso3   year diagnosis gender age   count population
##   <chr>      <chr> <chr> <dbl> <chr>      <chr> <chr> <dbl>      <dbl>
## 1 Afghanistan AF    AFG   1980 sp        m    014    NA        NA
## 2 Afghanistan AF    AFG   1980 sp        m   1524    NA        NA
## 3 Afghanistan AF    AFG   1980 sp        m   2534    NA        NA
## 4 Afghanistan AF    AFG   1980 sp        m   3544    NA        NA
## 5 Afghanistan AF    AFG   1980 sp        m   4554    NA        NA
## 6 Afghanistan AF    AFG   1980 sp        m   5564    NA        NA
```

- 3) Split the age column into two columns, min age and max age. Notice that there is no character separator. Check the documentation with `?separate` to understand other ways to separate the age column. Keep in mind that 0 to 14 is coded as 014 (3 characters) and the other age groups are coded with 4 characters. 65 only has two characters, but we will ignore that until the next problem.

```
tidy_who_pop <- tidy_who_pop %>%
  mutate(min_age = stringr::str_extract(age, "\\d{1,2}"),
         max_age = stringr::str_extract(age, "\\d{2}$"))

glimpse(tidy_who_pop)
```

```
## Rows: 405,440
## Columns: 11
## $ country    <chr> "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", ~
## $ iso2       <chr> "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", ~
## $ iso3       <chr> "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AF~
## $ year       <dbl> 1980, 1980, 1980, 1980, 1980, 1980, 1980, 1980, 1980, 1980, ~
## $ diagnosis  <chr> "sp", "sp", "sp", "sp", "sp", "sp", "sp", "sp", "sp", "sp", ~
## $ gender     <chr> "m", "m", "m", "m", "m", "m", "m", "f", "f", "f", "f", ~
## $ age        <chr> "014", "1524", "2534", "3544", "4554", "5564", "65", "014", ~
## $ count      <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ~
## $ population <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ~
## $ min_age    <chr> "01", "15", "25", "35", "45", "55", "65", "01", "15", "25", ~
## $ max_age    <chr> "14", "24", "34", "44", "54", "64", "65", "14", "24", "34", ~
```

- 4) Since we ignored the 65+ group in the previous problem we will fix it here. If you examine the data you will notice that 65 was placed into the max_age column and there is no value for min_age for those records. To fix this use mutate() in order to replace the blank value in the min_age column with the value from the max_age column and another mutate to replace the 65 in the max column with an Inf. Be sure to keep the variables as character vectors.

```
tidy_who_pop <- tidy_who_pop %>%
  mutate(min_age = ifelse(max_age == "65", max_age, min_age),
         max_age = ifelse(max_age == "65", "Inf", max_age))

glimpse(tidy_who_pop)
```

```
## Rows: 405,440
## Columns: 11
## $ country    <chr> "Afghanistan", "Afghanistan", "Afghanistan", "Afghanistan", ~
## $ iso2       <chr> "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", ~
## $ iso3       <chr> "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AF~
## $ year       <dbl> 1980, 1980, 1980, 1980, 1980, 1980, 1980, 1980, 1980, 1980, ~
## $ diagnosis  <chr> "sp", "sp", "sp", "sp", "sp", "sp", "sp", "sp", "sp", "sp", ~
## $ gender     <chr> "m", "m", "m", "m", "m", "m", "m", "f", "f", "f", "f", ~
## $ age        <chr> "014", "1524", "2534", "3544", "4554", "5564", "65", "014", ~
## $ count      <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ~
## $ population <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ~
## $ min_age    <chr> "01", "15", "25", "35", "45", "55", "65", "01", "15", "25", ~
## $ max_age    <chr> "14", "24", "34", "44", "54", "64", "Inf", "14", "24", "34"~
```

- 5) Find the count per diagnosis for males and females.

See ?sum for a hint on resolving NA values.

```
count_per_gender <- tidy_who_pop %>%
  group_by(gender, diagnosis) %>%
  summarise(count = n())
```


```
## 'summarise()' has grouped output by 'gender'. You can override using the
## '.groups' argument.
```

```
print(count_per_gender)
```

```
## # A tibble: 8 x 3
## # Groups:   gender [2]
##   gender diagnosis count
##   <chr>   <chr>     <int>
## 1 f      ep        50680
## 2 f      rel        50680
## 3 f      sn        50680
## 4 f      sp        50680
## 5 m      ep        50680
## 6 m      rel        50680
## 7 m      sn        50680
## 8 m      sp        50680
```

- 6) Now create a plot using ggplot and geom_col where your x axis is gender, your y axis represents the counts, and facet by diagnosis. Be sure to give your plot a title and resolve the axis labels.

```
ggplot(count_per_gender, aes(x = gender, y = count, fill = diagnosis)) +
  geom_col(position = "dodge") +
  facet_wrap(~ diagnosis, ncol = 3) +
  labs(title = "Counts per Gender by Diagnosis",
       x = "Gender",
       y = "Counts",
       fill = "Diagnosis") +
  theme_minimal()
```



7) Find the percentage of population by year, gender, and diagnosis. Be sure to remove rows containing NA values.

```
who_pop_percent <- tidy_who_pop %>%
  drop_na() %>%
  group_by(year) %>%
  mutate(total_pop = sum(population)) %>%
  group_by(year, gender, diagnosis) %>%
  summarise(percentage = sum(population)/first(total_pop))
```

'summarise()' has grouped output by 'year', 'gender'. You can override using
the '.groups' argument.

```
print(who_pop_percent)
```

```
## # A tibble: 94 x 4
## # Groups:   year, gender [38]
##   year gender diagnosis percentage
##   <dbl> <chr> <chr>      <dbl>
## 1 1995 f     sp         0.500
## 2 1995 m     sp         0.500
## 3 1996 f     sp         0.500
## 4 1996 m     sp         0.500
## 5 1997 f     sp         0.500
```

```
## 6 1997 m      sp      0.500
## 7 1998 f      sp      0.500
## 8 1998 m      sp      0.500
## 9 1999 f      ep      0.00117
## 10 1999 f     sn      0.00117
## # i 84 more rows
```

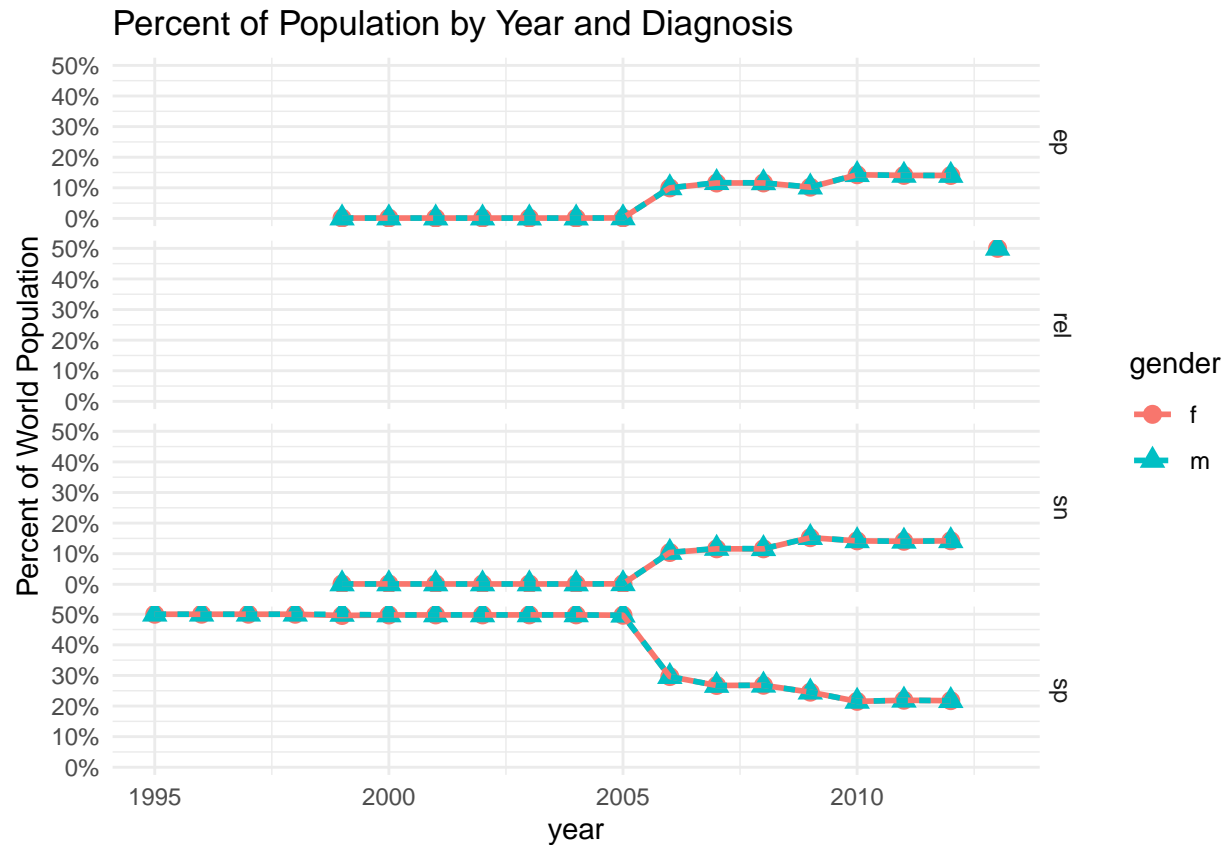
- 8) Create a line plot in ggplot where your x axis contains the year and y axis contains the percent of world population. Facet this plot by diagnosis with each plot stacked vertically. You should have a line for each gender within each facet. Be sure to format your y axis and give your plot a title.

```
who_pop_percent <- ungroup(who_pop_percent)

ggplot(who_pop_percent, aes(x = year, y = percentage * 100, color = gender, shape = gender, linetype = gender)) +
  geom_point(size = 3) +
  geom_line(size = 1) +
  facet_grid(diagnosis ~ .) +
  scale_y_continuous(labels = function(x) paste0(x, "%"), limits = c(0, NA)) +
  labs(title = "Percent of Population by Year and Diagnosis", y = "Percent of World Population") +
  theme_minimal()
```

```
## Warning: Using 'size' aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use 'linewidth' instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.
```

```
## 'geom_line()': Each group consists of only one observation.
## i Do you need to adjust the group aesthetic?
```



- 9) Now unite the min and max age variables into a new variable named `age_range`. Use a '-' as the separator.

```
tidy_who_pop <- tidy_who_pop %>%
  unite(age_range, min_age, max_age, sep = "-")

head(tidy_who_pop)
```

```
## # A tibble: 6 x 10
##   country iso2 iso3 year diagnosis gender age count population age_range
##   <chr>   <chr> <chr> <dbl> <chr>   <chr> <chr> <dbl>   <dbl> <chr>
## 1 Afghanist~ AF   AFG   1980 sp      m      014    NA      NA 01-14
## 2 Afghanist~ AF   AFG   1980 sp      m     1524    NA      NA 15-24
## 3 Afghanist~ AF   AFG   1980 sp      m     2534    NA      NA 25-34
## 4 Afghanist~ AF   AFG   1980 sp      m     3544    NA      NA 35-44
## 5 Afghanist~ AF   AFG   1980 sp      m     4554    NA      NA 45-54
## 6 Afghanist~ AF   AFG   1980 sp      m     5564    NA      NA 55-64
```

- 10) Find the percentage contribution of each age group by diagnosis. You will first need to find the count of all diagnoses then find the count of all diagnoses by age group. Join the former to the later and calculate the percent of each age group. Plot these as a `geom_col` where the x axis is the diagnosis, y axis is the percent of total, and faceted by age group.

```

# Count of all diagnoses
all_diagnoses_count <- tidy_who_pop %>%
  group_by(diagnosis) %>%
  summarise(total_count = n())

# Count of all diagnoses by age group and diagnosis
diagnoses_by_age_count <- tidy_who_pop %>%
  group_by(age_range, diagnosis) %>%
  summarise(age_diagnosis_count = n())

```

'summarise()' has grouped output by 'age_range'. You can override using the
'.groups' argument.

```

# Join the dfs
diagnosis_age_percent <- left_join(diagnoses_by_age_count, all_diagnoses_count, by = "diagnosis")

# Calculate the percentage of each age group by diagnosis
diagnosis_age_percent <- diagnosis_age_percent %>%
  mutate(percentage = (age_diagnosis_count / total_count) * 100)

# Plot the results
ggplot(diagnosis_age_percent, aes(x = diagnosis, y = percentage, fill = age_range)) +
  geom_col(position = "dodge") +
  facet_wrap(~ age_range, scales = "free") +
  labs(title = "Percent of Each Age Group by Diagnosis",
       x = "Diagnosis",
       y = "Percent of Total") +
  theme_minimal()

```


Percent of Each Age Group by Diagnosis

